Infants prefer to listen to speech: A meta-analysis.

Cécile Issard[1], Sho Tsuji[2], & Alejandrina Cristia[1]

[1] Laboratoire de Sciences Cognitives et Psycholinguistique, Ecole Normale Supérieure, Département d'Études Cognitives

[2] International Research Center for Neurointelligence, The University of Tokyo

Author Note

Correspondence concerning this article should be addressed to Cécile Issard, Laboratoire de Sciences Cognitives et Psycholinguistique, Département d'Études Cognitives, Ecole Normale Supérieure, 29 rue d'Ulm, 75005 Paris, France. E-mail: cecile.issard@gmail.com

11 Abstract

12

13 *Keywords:* Meta-analysis, infants, speech preference, auditory development, natural
14 sounds

15 Word count: X

Infants prefer to listen to speech: A meta-analysis.

Aknowledgement

This work was supported by an Agence Nationale de la Recherche grant to A.C. (ANR-17-CE28-0007 LangAge, ANR-16-DATA-0004 ACLEW, ANR-14-CE30-0003 MechELex, ANR-17-EURE-0017); and the J. S. McDonnell Foundation Understanding Human Cognition Scholar Award to A.C.

Introduction

Given the importance of speech for human vocal communication, it is conceivable that infants are born equipped with a preference for speech, which becomes stronger with age and exposure. Here, we synthesize empirical data on infants' preferences for speech over non-speech sounds to assess the explanatory role of three factors: stimulus naturalness, vocal quality, and familiarity.

[insert Figure 1 here]

Potential dimensions underlying preference patterns

There are three key conceptual explanations for infants preference for speech over competitor sounds: Preference for (a) natural over artificial sounds; (b) familiar over unfamiliar sounds; and (c) vocal over non-vocal sounds (Figure 1). These explanations are mutually compatible, and one or more may be true.

Natural versus artificial sounds. Natural sounds are those produced by biological systems, including vocal tracts but also the sound of walking and heart rate. In many cases natural and artificial sounds differ in their acoustic characteristics. For example,

backward speech has unnatural formant transitions and seemingly abrupt closures compared to naturally produced sounds. Natural sounds are processed more accurately by the auditory system, from the cochlea (Lewicki, 2006) to the auditory cortex (e.g. Mizrahi et al., 2014). This predicts a preference for speech over artificial competitors that is present from birth, consistent with existing literature (e.g., Vouloumanos & Werker, 2007).

Familiarity. Perhaps infants prefer speech to other sounds since it is a frequent sound. Newborns prefer their native speech to prosodically distinct foreign speech (e.g., Mehler et al., 1988), which supports a preference for sound patterns heard frequently in the womb. There are no behavioral results directly testing the prediction that infants show stronger preferences for speech over a foil when tested with more familiar speech stimuli (for instance, spoken in their native, as compared to a foreign language), but results from neuroimaging studies provide indirect evidence for this view. For instance, newborns' brain activation was different for forward than backward speech when the native language was used as the speech stimuli, but not when a foreign language was used (May et al., 2018).

Vocal versus non-vocal sounds. Many results summarized above may be accommodated by a third hypothesis, postulating a preference for vocal over non-vocal sounds. Vocal sounds are those made with a mouth, and thus typically a subset of natural sounds. Newborns made more head-turns to speech than to heartbeat (Ecklund-Flores & Turkewitz, 1996), arguably both equally natural and familiar to them, but they listened equivalently to speech and monkey calls, despite the greater familiarity of the former (Vouloumanos & Werker, 2010).

Changes as a function of development. Development may affect the preference for speech in various ways. Whereas newborns do not prefer speech over monkey calls, three-month-olds do (Vouloumanos et al., 2010). This suggests that as they age, infants might develop an increasingly narrow definition of the stimulus they prefer. In this case naturalness, familiarity, and vocal quality effects should change as a function of age: Very

close stimuli (e.g., speech versus another natural sound) initially leads to a weak preference, but, as infants age, this preference may be as strong as that found for very different stimuli (e.g., speech versus an artificial sound). Many articles discuss potential changes in the pattern of preference as a function of age (e.g. Ferry et al., 2013; Shultz & Vouloumanos, 2010; Shultz et al., 2014). To our knowledge, only two papers from the same laboratory include multiple age groups tested with the exact same stimulus categories and procedure (Vouloumanos et al., 2010; Vouloumanos & Werker, 2004). In fact, statements about age-related changes are often done using the demonstrably problematic method of concluding that there is an interaction without actually testing for it statistically (Gelman & Stern, 2012). It is therefore important to directly test these statements.

A meta-analytic approach

In sum, previous work on infants' preferences is broadly compatible with preference for natural over artificial, vocal over non-vocal, and familiar over unfamiliar sounds, potentially interacting with infants' age. In this paper, we seek to directly test these interpretations of the literature by employing a meta-analytic approach. Meta-analyses involve combining studies that may vary in their methodology. One limitation is therefore that one cannot isolate specific variables as well as in direct experimentation. Therefore, meta-analyses may miss subtle effects. Nonetheless, they have several useful features. They can reveal small effects not obvious in individual studies by combining them to obtain larger samples. Additionally, by integrating data across different laboratories, they provide evidence for the generalizability of effects across labs. Finally, meta-analyses offer tools to detect publication bias in the literature.

Specifically for the present case, a meta-analysis allows to statistically test different explanations. We can test the effect of factors that are not part of the original design, by redescribing the stimuli used as a function of those factors. For instance, a study measuring

preference for native speech over native backward speech provides data on a natural versus artificial, as well as a vocal versus non-vocal contrast. We can also draw a developmental timeline across the age range covered by the literature.

Meta-analyses can even provide theoretical and empirical insights that contradict qualitative reviews. For example, it has been proposed that infants' preference for novel or familiar items related to infants' age such that, all things equal, younger infants showed familiarity preferences whereas older infants exhibited novelty preferences (Hunter & Ames, 1988). However, Bergmann and Cristia (2016?) found stable familiarity preferences for word segmentation in natural speech across the first two years; and Black and Bergmann (2016) found a stable novelty effect for artificial grammars implemented in synthesized speech whereas those implemented in natural speech led to stable familiarity preferences. Meta-analyses are therefore important to statistically and systematically test the theoretical predictions proposed in qualitative reviews.

Given the scarcity of direct evidence on the potential explanations laid out above (naturalness, vocal quality, and familiarity, as a function of age), we conducted a meta-analysis to test whether infants' preference for speech sounds over other types of sounds is stable in newborns, and how it develops over the first year of life. Assuming all three factors are true, and further assuming that the definition of the preferred stimulus narrows with age, we predicted that infants will show (see Figure 2): 1. a greater preference for speech over natural sounds as a function of age, but a stable preference for speech over artificial sounds that is stable over development, 2. but an increasing preference for speech over natural sounds; 3. a greater preference for speech over other vocal sounds as a function of age, but a stable preference for speech over non-vocal sounds over development; 4. a greater preference for native speech over non-speech as a function of age, but a smaller preference for foreign speech over non-speech with age .

Methods

Literature search

We composed the initial list of studies with suggestions by experts (authors of this work); one google scholar searches (("speech preference" OR "own-species vocalization") AND infant - "infant-directed"), the same search in PubMed and PsycInfo (last searches on 24/09/2019); and a google alert. We also inspected the reference lists of all included papers.

Inclusion criteria

After a first screening based on titles and abstracts using more liberal inclusion criteria, we decided on inclusion based on full paper reading. We included studies that tested human infants from birth to 1 year of age, and contrasted speech sounds with any other type of sound, measuring behavioral responses to the sounds (e.g., looking times). We excluded studies that only contrasted foreign against native language, did not present natural speech sounds at all, presented speech in the mother's voice, or intentionally mixed speech with other vocal sounds within the same sound condition. We also excluded neuroimaging studies to avoid mixing results from different brain regions with different response profiles. We included published (i.e., journal articles) as well as unpublished works (i.e., doctoral dissertations as long as sufficient information was provided).

A PRISMA flow chart summarizes the literature review and selection process (Figure 3). We documented all the studies that we inspected in a decision spreadsheet (available in the online supplementary materials; Anonymized, 2019).

[Insert Figure 3 here]

Coding

The critical variables for our purpose are infant age, methodological variables (testing method: central fixation, high amplitude sucking, head-turn preference procedure, high amplitude sucking/passive listening), and key stimuli characteristics. Specifically, we coded the language in which the speech sounds were recorded (native or foreign), and whether the sound opposed to speech was natural or not, vocal or not. This competitor was coded as natural if it was produced by a biological organism without any further acoustic manipulation. If the authors applied acoustic manipulations it was coded as artificial. This sound was considered as vocal if it was produced by an animal vocal tract, either original or modified.

Data were coded by the first author. In addition, 20% of the papers were randomly selected to be coded by the second author independently, with disagreements resolved by discussion. There were 10 disagreements out of a total of 260 fields filled in indicative of the coders not following the codebook, which led to a revision of all data in four variables.

We coded all the statistical information reported in the included papers. If reported, we coded the mean score and the standard deviation for speech, and the other sound separately. When infant-level data was provided, we recomputed the respective mean scores and standard deviations based on the reported individual scores. If reported, we also coded the t-statistic between the two sound conditions, or an F-statistic provided this was a two-way comparison. If effect sizes were directly reported as a Cohen's d or a Hedges' g, we also coded this.

The PRISMA checklist, data, and code can be found on the online supplementary materials (Anonymized, 2019).

Effect sizes

Once the data were coded, we extracted effect sizes, along with their respective variance. Effect sizes were standardized differences (Cohen's d) between response to speech and to the other sound. If they were not directly reported in the papers, we computed them using the respective means and SDs ((**???**) & Wilson, 2001), or a t- or F-statistic ((**???**) et al., 1996). As our effect sizes came from within-subject comparisons (e.g. looking time of the same infant during speech and during monkey calls), we needed to take into account the correlation between the two measurements in effect sizes and effect size variances computations. We computed this correlation based on the t-statistic, the respective means and SDs ((**???**) & Wilson, 2001) if they were all reported; or imputed this correlation randomly if not. We finally calculated the variance of each effect size ((**???**) & Wilson, 2001). Cohen's d were transformed to Hedges' g by multiplying d by a correction for small sample sizes based on the degree of freedom (Borenstein et al., 2011). This procedure led to XX (some of them coming from the same infant group, hence not mutually independent) effect sizes, XX from published, peer-reviewed paper, and 1 from a thesis.

All analyses use the R [CITER] package Robumeta ((**???**) et al., 2010), which allows to fit meta-analytic regressions that take into account the correlated structure of the data, when repeated measures are obtained from the same infant groups within papers.

Results

Database description

We found a total of 28 papers reporting 47 (not mutually independent) effect sizes, see Table 1. 28 papers have been submitted to or published in peer-reviewed journals (citations). The remaining 1 paper was a thesis ((**???**)).

Studies tended to have small sample sizes, with a median N of 16 children (Range = 52, M = 20.66, Total: 659). Infants ranged from 0 to 12 months (1.50 to 380.50 days), although the majority were under 9 months of age (71.42% of the studies). Individual samples comprised 48 % of female participants on average. Infants were native of 7 different languages across the whole database (English, French, Japanese, Italian, Russian, Yiddish, Hebrew). Studies were performed in 13 different laboratories from 6 different countries (United States, Canada, Israel, France, Japan, Italy). 3 experimental methods were used: 12 studies used Central Fixation (CF) ((**???**) & Aslin,1994; (**???**) & Werker (2004); (**???**) et al. (2010); (**???**) & Vouloumanos (2010); (**???**) & Bundy (1981); (**???**) et al. (2009); (**???**) et al. (2019); (**???**) et al. (2019); (**???**) & Kishon-Rabin (2011); (**???**) & Vouloumanos (2013); (**???**) & Curtin (2014); (**???**), 2018); 3 used High-Amplitude Sucking (HAS) ((**???**) & Werker, 2007; (**???**) et al., 2010; (**???**) & DeCasper, 1987); and 1 used Head-turn Preference Procedure (HPP) ((**???**) & Turkewitz, 1996).

Summary effect size

Integrating across all studies in a meta-analytic regression without any moderator, we found a summary effect size g of 0.53 (SE = 0.09 CI = [0.34 , 0.73]), corresponding to a medium effect size.

Publication bias

We assessed the presence of a potential publication bias in the body of literature by studying the relationship between standard errors of effect sizes as a function of Hedges' g (see funnel plot in Figure 4). A regression test on these data was significant ($z = 6.40$, $p < 0.0001$), as was the Kendall's tau rank correlation test for funnel plot asymmetry (Kendall's tau = 0.5172, $p < .0001$), indicating a publication bias in the literature. To further

investigate this bias, we symmetrized the funnel plot with the "trim and fill" method (Duval, 2005). XX18 (SE = XX5.87) missing studies were needed on the left side of the plot to symmetrize the funnel plot.

[insert Figure 4 here]

Moderator analyses

We then tested if the preference found above could be explained by the dimensions discussed in the literature. Following our hypothesis, we fit a meta-analytic model with the following moderators:

- mean age of children;
- familiarity with the language used (native or foreign);
- naturalness of the contrastive sound (coded as yes if it was natural and no otherwise).
- vocal quality of the contrastive sound (coded as yes if it was vocal and no otherwise).

There was only a significant effect of the experimental method used (central fixation led to higher effect sizes than the other methods; see Figure 5). None of the other moderators was significant (see Table 1).

Due to the relatively low number of effect sizes available in the literature, we did not add interactions with age to the model to avoid overfitting. Inspection of results in Figure 5 show that the confidence intervals of all conditions overlap almost exactly across the tested ages, excluding the possibility of such interactions.

Readers may wonder to what extent our results are obscured by the influence of experimental method, which is known to be sizable in the cognitive developmental literature ((**???**) et al., 2018). Inspection of plots where effect sizes has been residualized from experimental method (Supplementary figure SX) looks virtually identical.

Discussion

Our meta-analysis synthesizes the available literature on infants' preference for speech sounds. Our results confirm that infants reliably prefer speech over other types of sounds from birth. When all studies were considered together with no moderators, we found a sizable intercept (g=.XX). For comparison, the main effect for native vowel discrimination using looking time methods is estimated at .25 (Tsuji & Cristia, 2013; data inspected in metalab.stanford.edu on 2019-10-18). We had predicted infants' speech preference to be larger when the competitor was an artificial sound than when it was a natural one; when the competitor was non-vocal; and when the speech was in the infants' native language. In fact, we were unable to disprove the null hypothesis of no difference for all three factors, with widely overlapping distributions of effect sizes for studies varying along the three dimensions.

We had also hypothesized age to play a major role, because it may correlate with a reshaping of the category definition for speech itself. Indeed, studies comparing processing of human speech against human non-speech as well as animal vocalizations more generally (Vouloumanos et al., 2010; Mac Donald et al., 2019) often discuss these age-related differences in categorization of these sounds. Surprisingly, age did not significantly moderate the overall preference for speech.

From birth and regardless of changes co-occurring with age, infants show a preference for speech, which cannot be reduced to three simpler explanations: naturalness, vocalness, or familiarity (represented here by the native/foreign contrast). This capacity to preferentially listen to speech sounds from birth suggests that infants are born with the capacity to recognize their conspecifics' communication signals. This parallels what has been proposed by the Conspec model for faces: infants would be born with knowledge about faces, enabling them to orient their attention toward them, even without any prior exposure to faces (Morton & Jonhson, 1991). The fact that familiarity with the language used in the

experiment did not modulate infants' preference suggest that exposure did not play a crucial role for speech either. However, contrary to faces, fetuses are exposed to speech that is low-pass filtered by the womb throughout the last trimester of gestation (Querleu et al., 1988, Lecanuet & Granier-Deferre, 1993). It is therefore possible that prenatal experience with low-pass filtered speech helps infants to form a representation of speech, independently of the language spoken.

The Conspec model proposes that faces would be detected because of their spatial structure (Morton & Johnson, 1991). Similarly, it is possible that infants prefer speech because of its complex acoustic structure and fast transitions (Rosen, 2007). Speech is characterized by joint spectral and temporal modulations at specific rates (Singh & Theunissen, 2003). It is possible that infants attune to this specific spectro-temporal structure (though see **???**, for evidence that this explanation alone may not be sufficient for neural responses). Testing this explanation would require to carry out acoustic analyses of the actual stimuli used in the studies. Thus, we recommend interested researchers to gather more data in which the competitor is acoustically simple versus complex; and to deposit the actual stimuli in a public archive such as the Open Science Framework (REF).

One may wonder whether we fail to find many differences because of a lack of statistical power, particularly in view of between-study variability. We think it is unlikely that all null results reported in this paper are due to this. Inspection of results in Figure 3 show that the confidence intervals of all conditions overlap almost exactly. Moreover, Table XX shows that the estimate for all these factors is close to zero (the maximum being XX). That said, more data would be welcome to confirm our results with more statistical power. It would also be important to carry out more tests on infants older than 9 months. Language production gains in complexity at about this age [(**???**), which could affect infants' speech preference. We particularly recommend using as competitor natural vocal stimuli, and as target foreign speech, which would help fill in an important gap in our dataset.

Another finding of our meta-analysis is that the distribution of effect sizes in the literature is consistent with publication bias, in view of a strong asymmetry of the funnel plot. In fact, the trim-and-fill method suggested XX points may be missing, which is a considerable number given that we have XX effect sizes in total (i.e., a quarter more would be missing). Unsurprisingly, the missing studies are in the negative section, i.e., a preference *against* speech, a result that could lead authors to doubt their own data and not submit it to journals, or that would be considered odd by reviewers and editors, who may ask that the data be removed (or who may recommend the paper to be rejected altogether). These missing studies constitute an important limitation of our results. The literature being biased toward positive effect sizes, the true effect size might be smaller than the one we found (vertical line on Figure X).

Ultimately, preferential processing of speech may support higher level cognitive tasks. The human species is a highly social one. Detecting speech signals would allow to integrate it with other sensory percepts, such as faces, to form multisensory representations of conspecifics ((**???**) et al., 2009). This would lay the track for social cognition. Identifying speech signals and paying attention to them would allow infants to form complex representations of the sensory world, that they can manipulate cognitively. Infants could categorize visual stimuli (i.e. associate a label to a category of objects) when they were associated to speech, but not pure tones or backward speech ((**???**) & Waxman, 2007; (**???**) et al, 2010; (**???**) et al, 2013). Interestingly, infants categorized visual stimuli when presented with speech, melodies, or monkey vocalizations (Fulkerson & Haaf, 2003; Ferry et al, 2013). These results support the idea that infants may preferentially process complex sounds. Finally, the preference itself may also be a meaningful index of processing that can be used to identify children at risk ((**???**) et al., 2019). It is therefore important to take stock of what we know today.

Given the crucial importance of understanding infants' speech preference, we make the

following recommendations for further data collection, analysis, and reporting. First, authors should strive to increase their sample sizes. The median sample size at present is 20, which is close to the field standard ((**???**) et al., 2018) but much lower than current recommendations ((**???**), 2017). Second, authors should consider proposing their studies as registered reports ((**???**)). In this new publication scheme (available for Developmental Science, Infancy, Infant Behavior and Development, and Journal of Child Language at the time of writing, see a full up-to-date list on https://cos.io/rr/), manuscripts are submitted before data are collected. Reviewers and editors make publication decisions based solely on the introduction and methods. Once the paper is accepted, the author collects the data, analyses it according to a pipeline described in the accepted methods, and writes up the rest of the manuscript. The paper is then reviewed once more for readability, but it cannot be rejected if the results are surprising or uncomfortable for the field. Third, for authors who would rather not follow this publication route, we still strongly recommend the use of pre-specified analysis plans. These have the virtue of, when used correctly, allowing both authors and readers to separate confirmation from exploration, reducing the likelihood of inadvertently engaging in questionable research practices (**???**), known to increase false positives. Finally, we strongly recommend reviewers and editors to evaluate submitted manuscripts on the basis of the quality of the methods, and not of the results. If a study fails to report a speech preference, or actually reports a preference for the competitor, this may actually reflect the reality of this phenomenon. For interested readers who intend to collect such data, we recommend caution when designing the study, and transparency when reporting it. Our dataset can be community-augmented, and we invite researchers investigating this phenomenon to complement it with any data they would have (Anonymized link), whatever the results and publication status.